

Lab 04

Profiling, debugging and class templates

Paolo Joseph Baioni

October 31, 2025

Recalls

Introduction to exercise 1

- ▶ The implemented matrix class is organized as **column-major**, *i.e.* $A(i, j) = \text{data}[i + j * \text{rows()}]$, conversion from 1d to 2d indexing is performed by the utility method `sub2ind`.
- ▶ Access to elements is implemented both in `const` and non-`const` versions, by overloading `operator()`.
- ▶ Data is private, *getter methods* expose what is needed to the user, both `const` and non-`const` versions are provided.
- ▶ Naive implementation of matrix-matrix multiplication is slow because it has low *data locality*, simply transposing the left matrix factor improves performance significantly¹.
- ▶ The `#include <ctime>` header provides timing utilities, `tic()` and `toc(x)` macros start and stop the timer.

¹See M. Kowarschik, C. Weiß. (2002). *Lecture Notes in Computer Science*. 213-232. DOI: 10.1007/3-540-36574-5_10 for further details.

Follow-up of Lab 3

Exercise 1.3

- ▶ Include the `Eigen/Dense` header.
- ▶ Use the `Eigen::Map` template class to wrap the matrix data and interpret it as `Eigen::MatrixXd`.
- ▶ Compare the execution speed with respect to the previous implementations for varying matrix size, such as `msize=500` and `msize=5000`.

Profiling and memory checking

Exercise 1.4

Going back to Ex. 1.2, perform the subsequent analysis:

- ▶ `coverage (lcov)`
- ▶ `memcheck (valgrind)`
- ▶ `profile (valgrind, kcachegrind)`

Debugging

Exercise 2

The program `integer-list` in the directory `02-bug` has:

- ▶ a compile error;
- ▶ a run-time error;
- ▶ a memory leak;
- ▶ a potential memory leak that is not captured by the `main`.

Find all the issues and fix them.

Get help by using `gdb` and `valgrind`.

Class templates

Finite differences

Implement a C++ template class to evaluate derivatives of any order of a given function (callable object) at a given point using recursive backward/forward first-order finite difference formulas.

Class templates

Sparse matrix introduction

- ▶ A $N \times N$ sparse matrix is a matrix whose number of non-zero elements N_{nz} is $O(N)$.
- ▶ The average number m of non-zero elements per row (or column) is constant w.r.to the matrix size.
- ▶ If the majority of matrix entries is 0, *i.e.* if $m \ll N$ it is convenient to store only the non-zero elements.
- ▶ The matrix-vector product (which is the basic ingredient of Krylov solvers) costs $O(N)$ rather than $O(N^2)$.

Class templates

Sparse matrix formats

Some (slightly revisited) classical data structures for sparse matrices

$$A = \begin{bmatrix} 4 & -1 & 0 & 0 \\ -1 & 4 & -1 & 0 \\ 0 & -1 & 4 & -1 \\ 0 & 0 & -1 & 4 \end{bmatrix}$$

COO (coordinates) or AIJ:

```
std::vector<double>      A{4, -1, -1, ...};  
std::vector<unsigned int> I{0, 0, 1, ...};  
std::vector<unsigned int> J{0, 1, 0, ...};
```

triplet vector:

```
std::vector<std::tuple<unsigned int, unsigned int, double>>  
    t{{0, 0, 4}, {0, 1, -1}, {1, 0, -1}, ...};
```

CSR (Compressed Sparse Row) or CRS (Compr. Row Storage) or Yale:

```
std::vector<double>      val{4, -1, -1, 4, -1, -1, 4, -1, -1, 4};  
std::vector<unsigned int> col_idx{0, 1, 0, 1, 2, 1, 2, 3, 2, 3};  
std::vector<unsigned int> row_ptr{0, 2, 5, 8, 10}; // n_rows + 1.
```


Class templates

Sparse matrix typical operations

- Insertion:

```
A[i][j] = x;
```

- Deletion:

```
A[i].erase(j); or A.erase(i, j);
```

- Random access:

```
x = A[i][j]; A[i][j] += y;
```

- Sequential traversing:

```
for (row : A)
{
    for (column : row)
        std::cout << column.value << " ";
    std::cout << std::endl;
}
```

- Matrix-vector multiplication:

```
std::vector<double> y = A * x;
```

Class templates

Inheriting from STL containers

The C++ standard is very permissive for the implementation of new containers, but:

- ▶ STL containers have non-virtual destructors!
- ▶ C.35: A base class destructor should be either public and virtual, or protected and non-virtual.

```
class Base {  
public:  
    ~Base { do_something(); }; // Non-virtual.  
}
```

```
class Derived : public Base {  
public:  
    MyComplexType member;  
    ~Derived { member.clear(); ... }  
}
```

```
Base *var = new Derived;  
delete var; // Calls var::~~Base() but not var::~~Derived()!
```

Class templates

Sparse matrix exercise

- ▶ Implement a C++ class to represent a sparse matrix inheriting from suitable STL containers.
- ▶ Simplify random access, allocation, entry increment, sequential traversing.
- ▶ (Optional) Refactor the code: instead of inheriting, make the actual STL container a private class member. Which operators should be added for the code to compile?

Class templates

Sparse matrix exercise

Implement the following methods:

```
/// Convert row-oriented sparse matrix to AIJ format.
```

```
void
```

```
aij(std::vector<double> &    a,  
    std::vector<unsigned int> &i,  
    std::vector<unsigned int> &j);
```

```
/// Convert row-oriented sparse matrix to CSR format.
```

```
void csr(std::vector<double> &    a,  
         std::vector<unsigned int> &col_ind,  
         std::vector<unsigned int> &row_ptr);
```

```
/// Stream operator.
```

```
friend std::ostream &
```

```
operator<<(std::ostream &stream, Class templates){Sparse matrix exercise}  
          sparse_matrix &M);
```

```
/// Sparse matrix increment.
```

```
void sparse_matrix::operator+=(sparse_matrix &other);
```

```
/// Compute matrix-vector product.
```

```
friend std::vector<double>
```

```
operator*(sparse_matrix &M,  
          const std::vector<double> &x);
```